

Introduction to Machine Learning

Individual Laboration Report –6–

Erik Sven Vasconcelos Jansson
erija578@student.liu.se
Linköping University, Sweden

December 15, 2016

Finally, the last machine learning topic covered are *artificial neural networks*. These estimators are very flexible, such that even a *single layer feed-forward neural network* complies with the *universal approximation theorem*, presented by Csáji [Csá01]:

Theorem 1 (Universal Approximation Theorem). *Any artificial feed-forward neural network with a single hidden layer, containing a finite amount of neurons, can approximate any continuous functions on the compact subset \mathcal{R}^n (with restrictions on σ).*

Proof. Csáji's [Csá01] derivation of Theorem 1. \square

Basically, the theorem states that even simple neural networks can represent interesting functions, given some suitable subset of *activation functions*. For this assignment, we want to approximate $\sin x$, where we are given 25 observation for training set. Also, we are given a validation set of length 25 for checking if our neural network is under/overfitting. We are using the *R* package *neuralnet* for our fit with 10 hidden units in a single hidden layer, also initialized with random weights in $[-1, 1]$ interval. See Listing 1 for the entire assignment source code.

For the curious, Equations 1, 2, and 3 are given:

$$\sigma(u) = \frac{1}{1 + e^{-u}} \quad (1)$$

$$\mathbf{w}_{(i)} = \mathbf{w}_{(i-1)} - \eta_k \nabla E(\mathbf{w}_{(i-1)}) \quad (2)$$

$$\hat{y}_j(\mathbf{x}) = \sigma(w_0 + \sum_{h=1}^H \sigma(w_{0h} + \mathbf{w}_h^T \mathbf{x})) \quad (3)$$

1. **Sigmoid Activation Function:** “S”-shaped function which converges $\sigma(u) = 1$ as $u \rightarrow \infty$ and $\sigma(u) = 0$ as $u \rightarrow -\infty$. Used in Equation 3.
2. **Batch Gradient Descent:** finds the “step” in the right direction for minimizing error E . This is achieved with the *gradient* of E given in respect to the weights \mathbf{w} ; giving a *hyperplane*.
3. **Single-Layer Neural Network Estimator:** uses Equations 1 and 2 to find \hat{y}_j by finding the parameters \mathbf{w} in each layer (a linear equation) by means of *gradient descent* and producing a *non-linear result in subsequent layers* by the *activation function*. This is the primary reason why neural networks are so flexible & general.

By using a *threshold* for the *gradient descent* we can stop the *neural network* from either *overfitting* or *underfitting*. This simply done by increasing the threshold iteratively and taking the validation set's:

Threshold	S.S.E.
0.001	0.01367691527
0.002	0.01262419958
0.003	0.00988418900
0.004	0.00850089424
0.005	0.00955545744
0.006	0.00974372099
0.007	0.01583926857
0.008	0.01649252416
0.009	0.02112490377
0.010	0.02735909554

Table 1: Neural Network Values

After finding the “optimal threshold” of 0.004 by picking the 4^{th} iteration (where $i = 4$ that is) which gives the least amount of error for a validation set, we plot the best neural network in Figure 1, and also the predictions in Figure 2 for a *sine function*. Notice that the fit is pretty good, and the estimator gives a pretty “spot on” prediction for the function. It seems *neural networks* are incredibly powerful, but take time to train and are harder to reason about (for example, how do we choose the number of hidden layers and units? How long will it take?)

References

- [Csá01] Balázs Csanád Csáji. Approximation with Artificial Neural Networks. *Faculty of Sciences, Eötvös Loránd University, Hungary*, 24:48, 2001.
- [FHT09] Jerome Friedman, Trevor Hastie, and Robert Tibshirani. *The Elements of Statistical Learning*. Springer series in statistics, Berlin, second (11th) edition, 2009.
- [GF10] Frauke Günther and Stefan Fritsch. neuralnet: Training of Neural Networks. *The R Journal*, 2(1):30–38, 2010.

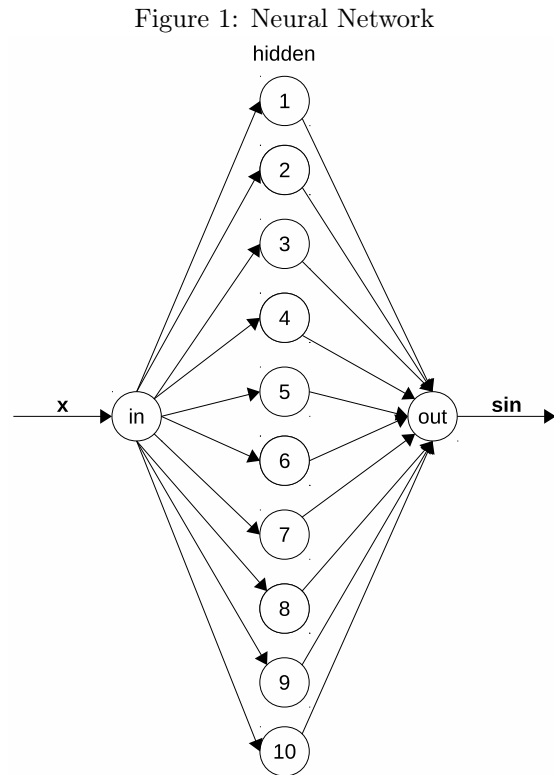



Figure 2: Neural Network’s Produced Predictions (in the graph are **raw values** and **predicted values**).



Appendix

Listing 1: Feed-Forward Backpropagating Neural Network Sine Estimator Script

```
1 library("ggplot2")
2 library("reshape2")
3 library("neuralnet")
4 library("grDevices")
5 set.seed(1234567890)
6
7 variable <- runif(50, 0, 10)
8 sine <- data.frame(x=variable, sin=sin(variable))
9 training <- sine[1:25,] ; testing <- sine[26:50,]
10
11 candidate_error <- Inf
12 units <- 10 # Hidden baby!
13 candidate_threshold <- Inf
14 weights <- runif(50, -1, +1)
15
16 for (threshold_attempt in 1:10) {
17   thresholdi <- threshold_attempt / 1000
18   nn <- neuralnet(sin~x, training, units,
19                 startweights = weights,
20                 threshold = thresholdi)
21
22   predicted <- compute(nn, testing$x)
23   error <- sum((testing$sin - predicted$net.result)^2)
24   cat("NN Threshold", thresholdi, "->", error, "SSE \n")
25   if (error < candidate_error) {
26     candidate_error = error
27     candidate_threshold = thresholdi
28   }
29 }
30
31 nn <- neuralnet(sin~x, training, units,
32               candidate_threshold,
33               startweights = weights)
34 predicted <- compute(nn, testing$x)
35
36 plot(nn)
37 setEPS()
38 cairo_ps("predictions.eps")
39 plot(testing$x, predicted$net.result, col = "red",
40      xlab = "x", ylab = "sin")
41 points(sine, col = "blue")
42 dev.off()
```

Listing 2: Output About the Produced Neural Network in the Assignment

```
1 $response
2      sin
3 1  0.31115890803
4 2 -0.65787112371
5 3  0.85356988285
6 4  0.92820698816
7 5  0.71194544538
8 6  0.95969186755
9 7  0.27531467859
10 8 -0.03662256168
```

```

11 9 -0.29718457265
12 10 -0.43427724087
13 11 0.27176755816
14 12 0.96762993527
15 13 0.87023024548
16 14 0.90319426880
17 15 0.53211225475
18 16 -0.90515370065
19 17 -0.99209419164
20 18 0.75493516282
21 19 0.43639270658
22 20 0.42400734122
23 21 0.77254200174
24 22 0.68138797265
25 23 0.32070401674
26 24 -0.99484612705
27 25 -0.82027249428
28
29 $covariate
30      [,1]
31 [1,] 9.1083657276
32 [2,] 3.8595812093
33 [3,] 8.4019783861
34 [4,] 7.4727497180
35 [5,] 7.0754500036
36 [6,] 7.5690891198
37 [7,] 0.2789170155
38 [8,] 9.4614087138
39 [9,] 9.7265206091
40 [10,] 9.8740137112
41 [11,] 9.1495487187
42 [12,] 1.3156641065
43 [13,] 1.0556694935
44 [14,] 7.4103393406
45 [15,] 6.8442786904
46 [16,] 4.2733338126
47 [17,] 4.5865617390
48 [18,] 8.5692227143
49 [19,] 2.6900070859
50 [20,] 0.4378655413
51 [21,] 0.8828348410
52 [22,] 7.0328426105
53 [23,] 2.8151199827
54 [24,] 4.8139597056
55 [25,] 4.1034799209
56
57 $err.fct
58 function (x, y)
59 {
60     1/2 * (y - x)^2
61 }
62 <environment: 0x339a758>
63 attr("type")
64 [1] "sse"
65
66 $act.fct
67 function (x)
68 {
69     1/(1 + exp(-x))
70 }
71 <environment: 0x339a758>
72 attr("type")

```

```

73 [1] "logistic"
74
75 $linear.output
76 [1] TRUE
77
78 $data
79      x      sin
80 1  9.1083657276  0.31115890803
81 2  3.8595812093 -0.65787112371
82 3  8.4019783861  0.85356988285
83 4  7.4727497180  0.92820698816
84 5  7.0754500036  0.71194544538
85 6  7.5690891198  0.95969186755
86 7  0.2789170155  0.27531467859
87 8  9.4614087138 -0.03662256168
88 9  9.7265206091 -0.29718457265
89 10 9.8740137112 -0.43427724087
90 11 9.1495487187  0.27176755816
91 12 1.3156641065  0.96762993527
92 13 1.0556694935  0.87023024548
93 14 7.4103393406  0.90319426880
94 15 6.8442786904  0.53211225475
95 16 4.2733338126 -0.90515370065
96 17 4.5865617390 -0.99209419164
97 18 8.5692227143  0.75493516282
98 19 2.6900070859  0.43639270658
99 20 0.4378655413  0.42400734122
100 21 0.8828348410  0.77254200174
101 22 7.0328426105  0.68138797265
102 23 2.8151199827  0.32070401674
103 24 4.8139597056 -0.99484612705
104 25 4.1034799209 -0.82027249428
105
106 $net.result
107 $net.result[[1]]
108      [,1]
109 1  0.30018554412
110 2 -0.64466078637
111 3  0.82577426828
112 4  0.95531707389
113 5  0.71041413678
114 6  0.98604140527
115 7  0.27220011825
116 8 -0.02393016651
117 9 -0.27977135370
118 10 -0.42452441939
119 11  0.26376340300
120 12  0.97423750701
121 13  0.86489547875
122 14  0.92926095080
123 15  0.49438343511
124 16 -0.91926885456
125 17 -0.98908629285
126 18  0.72304892189
127 19  0.42816650637
128 20  0.43013044493
129 21  0.76891173652
130 22  0.67393025023
131 23  0.32904062832
132 24 -0.97934291894
133 25 -0.83193245662
134

```

```

135
136 $weights
137 $weights[[1]]
138 $weights[[1]][[1]]
139           [,1]           [,2]           [,3]           [,4]           [,5]
140 [1,] 0.3718763846 -10.931812117  8.275060257  7.8216380497  1.551228805
141 [2,] 0.5081317757  1.628735531 -2.289303563  0.1166254588 -0.594052483
142           [,6]           [,7]           [,8]           [,9]           [,10]
143 [1,]  4.7453831203 -0.6070429987  9.38110372186 -0.1377222063  5.958245683
144 [2,] -0.4968713811  0.1924524647  0.09270470267  3.1685937697 -2.602280174
145
146 $weights[[1]][[2]]
147           [,1]
148 [1,] -0.06529714022
149 [2,] -0.70033511720
150 [3,]  3.84669442716
151 [4,]  2.74586539382
152 [5,] -0.75978348453
153 [6,] -9.09090264177
154 [7,]  6.65416477397
155 [8,] -8.24869453812
156 [9,] -0.20335986240
157 [10,] 1.00884789102
158 [11,] 2.01492912933
159
160
161
162 $startweights
163 $startweights[[1]]
164 $startweights[[1]][[1]]
165           [,1]           [,2]           [,3]           [,4]           [,5]
166 [1,] 0.4591262657 -0.5031667114  0.9014065554  0.9589186665  0.3389983536
167 [2,] 0.5853618421  0.4307389595 -0.8788680169  0.4978831881 -0.5358421607
168           [,6]           [,7]           [,8]           [,9]           [,10]
169 [1,]  0.6293357364 -0.4028865024  0.5968314419 -0.07648990629  0.7421438890
170 [2,] -0.4464168334 -0.3094406570  0.9041418806 -0.17025629710 -0.8588890089
171
172 $startweights[[1]][[2]]
173           [,1]
174 [1,] -0.2698646844
175 [2,] -0.9254528601
176 [3,]  0.3498687637
177 [4,] -0.7230960354
178 [5,] -0.9836924160
179 [6,] -0.6452815034
180 [7,]  0.8491520174
181 [8,]  0.6410233583
182 [9,] -0.4020887311
183 [10,] 0.9406797229
184 [11,] 0.2142777084
185
186
187
188 $generalized.weights
189 $generalized.weights[[1]]
190           [,1]
191 1  -4.18437409317
192 2   0.84720544205
193 3  -3.90202008113
194 4   8.87306916010
195 5   4.06386829427
196 6  18.83336456829

```

```

197 7 5.28857232424
198 8 38.75962317972
199 9 2.72921588482
200 10 1.62852618574
201 11 -4.58128912031
202 12 13.07325336492
203 13 4.31497701909
204 14 6.93919826366
205 15 4.07243773606
206 16 0.23010347764
207 17 0.02853630559
208 18 -3.31477743585
209 19 -3.26159533671
210 20 3.80902182260
211 21 3.41516125086
212 22 3.98677345674
213 23 -3.57280260613
214 24 -0.06878886818
215 25 0.40865817424
216
217
218 $result.matrix
219
220 error 1
221 reached.threshold 0.003576080337
222 steps 23174.000000000000
223 Intercept.to.l1ayhid1 0.371876384634
224 x.to.l1ayhid1 0.508131775660
225 Intercept.to.l1ayhid2 -10.931812117300
226 x.to.l1ayhid2 1.628735530657
227 Intercept.to.l1ayhid3 8.275060257474
228 x.to.l1ayhid3 -2.289303563425
229 Intercept.to.l1ayhid4 7.821638049688
230 x.to.l1ayhid4 0.116625458773
231 Intercept.to.l1ayhid5 1.551228805249
232 x.to.l1ayhid5 -0.594052483023
233 Intercept.to.l1ayhid6 4.745383120333
234 x.to.l1ayhid6 -0.496871381057
235 Intercept.to.l1ayhid7 -0.607042998673
236 x.to.l1ayhid7 0.192452464714
237 Intercept.to.l1ayhid8 9.381103721859
238 x.to.l1ayhid8 0.092704702673
239 Intercept.to.l1ayhid9 -0.137722206303
240 x.to.l1ayhid9 3.168593769723
241 Intercept.to.l1ayhid10 5.958245682591
242 x.to.l1ayhid10 -2.602280174422
243 Intercept.to.sin -0.065297140222
244 l1ayhid.1.to.sin -0.700335117198
245 l1ayhid.2.to.sin 3.846694427157
246 l1ayhid.3.to.sin 2.745865393824
247 l1ayhid.4.to.sin -0.759783484527
248 l1ayhid.5.to.sin -9.090902641770
249 l1ayhid.6.to.sin 6.654164773966
250 l1ayhid.7.to.sin -8.248694538124
251 l1ayhid.8.to.sin -0.203359862396
252 l1ayhid.9.to.sin 1.008847891021
253 l1ayhid.10.to.sin 2.014929129330
254
255 attr(,"class")
256 [1] "nn"

```
